

Orthogonal Optimization Of Subqueries And Aggregation

February 13, 2025

1 Introduction

In this paper, we present subquery and aggregation techniques implemented in Microsoft SQL Server.

1.1 Standard subquery execution strategies

Before describing subquery strategies in detail, it is important to clarify the two forms of aggregation in SQL, whose behavior diverges on an empty input.

“Vector” aggregation specifies grouping columns as well as aggregates to compute.

```
1 select o_orderdate, sum(o_totalprice)
2 from orders
3 group by o_orderdate
```

And there are queries that *always returns exactly one row*:

```
1 select sum(o_totalprice) from orders
```

In algebraic expressions we denote vector aggregate as $\mathcal{G}_{A,F}$, where A are the grouping columns and F are the aggregates to compute; and denote scalar aggregate as \mathcal{G}_F^1

We review standard subquery execution strategies using the following SQL query, which finds customers who have ordered more than \$1,000,000.

```

1  -- Q1
2  select c_custkey
3  from customer
4  where 100000 <
5         (select sum(o_totalprice)
6          from orders
7          where o_custkey = c_custkey)

```

Outerjoin, then aggregate:

```

1  select c_custkey
2  from customer left outer join
3         orders on o_custkey = c_custkey
4  group by c_custkey
5  having 1000000 < sum(o_totalprice)

```

Aggregate, then join:

```

1  select c_custkey
2  from customer,
3       (select o_custkey from orders
4        group by c_custkey
5        having 1000000 < sum(o_totalprice))
6  as aggresult
7  where o_custkey = c_custkey

```

1.2 Our technique: Use primitive, orthogonal pieces

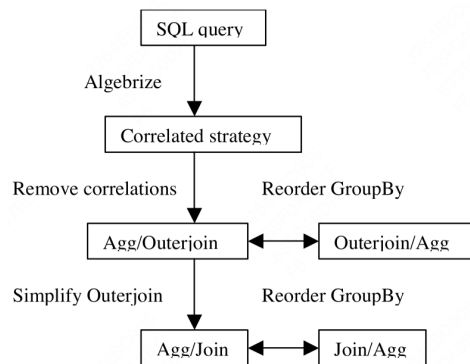


Figure 1: Primitives connecting different execution strategies

By implementing all these orthogonal techniques, the query processor should then produce the same efficient execution plan for the various equivalent SQL formulations we have listed above, achieving a degree of syntax-independence.

- **Algebrize into initial operator tree**
- **Remove correlations**
- **Simplify outerjoin**
- **Reorder GroupBy**

1.3 A useful tool: Represent parameterized execution algebraically

Apply takes a relational input R and a parameterized expression $E(r)$; it evaluates expression E for each row $r \in R$, and collects the results. Formally,

$$R\mathcal{A}^{\otimes}E = \bigcup_{r \in R} (\{r\} \otimes E(r))$$

where \otimes is either cross product, left outerjoin, left semijoin, or left antijoin.

The most primitive form is \mathcal{A}^{\times} , and cross product is assumed if no join variant is specified.

All operators used in this paper are bag-oriented, and we assume no automatic removal of duplicates. In particular, the union operator above is UNION ALL. Duplicates are removed explicitly using DISTINCT.

Now Q1 is like

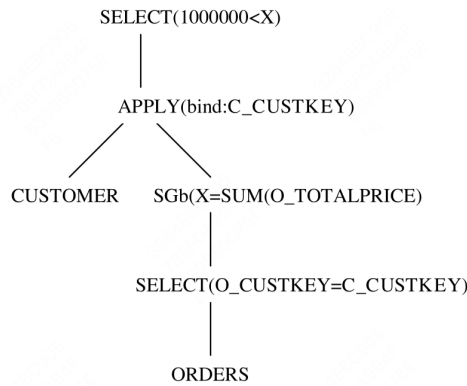


Figure 2: Subquery execution using Apply

Apply works on expressions that take scalar (or row-valued) parameters. A second useful construct is **SegmentApply**, which deals with expressions using *table*-valued parameters. It takes a relation input R , a parameterized expression $E(S)$, and a set of segmenting columns A from R . It creates segments of R using columns A , much like GroupBy, and for each such segment S it executes $E(S)$. Formally,

$$R \mathcal{S} \mathcal{A}_A E = \bigcup_a (\{a\} \times E(\sigma_{A=a} R))$$

where a takes all values in the domain of A .

2 Representing and normalizing subqueries

2.1 Direct algebraic representation with mutual recursion

2.2 Algebraic representation with Apply

2.3 Removal of Apply

$$R \mathcal{A}^\otimes E = R \otimes_{\text{true}} E \quad (1)$$

if no parameters in E resolved from R

$$R \mathcal{A}^\otimes (\sigma_p E) = R \otimes_p E \quad (2)$$

if no parameters in E resolved from R

3 Problems

4 References

References